

# Intermediate RStudio Training

January 23, 2018



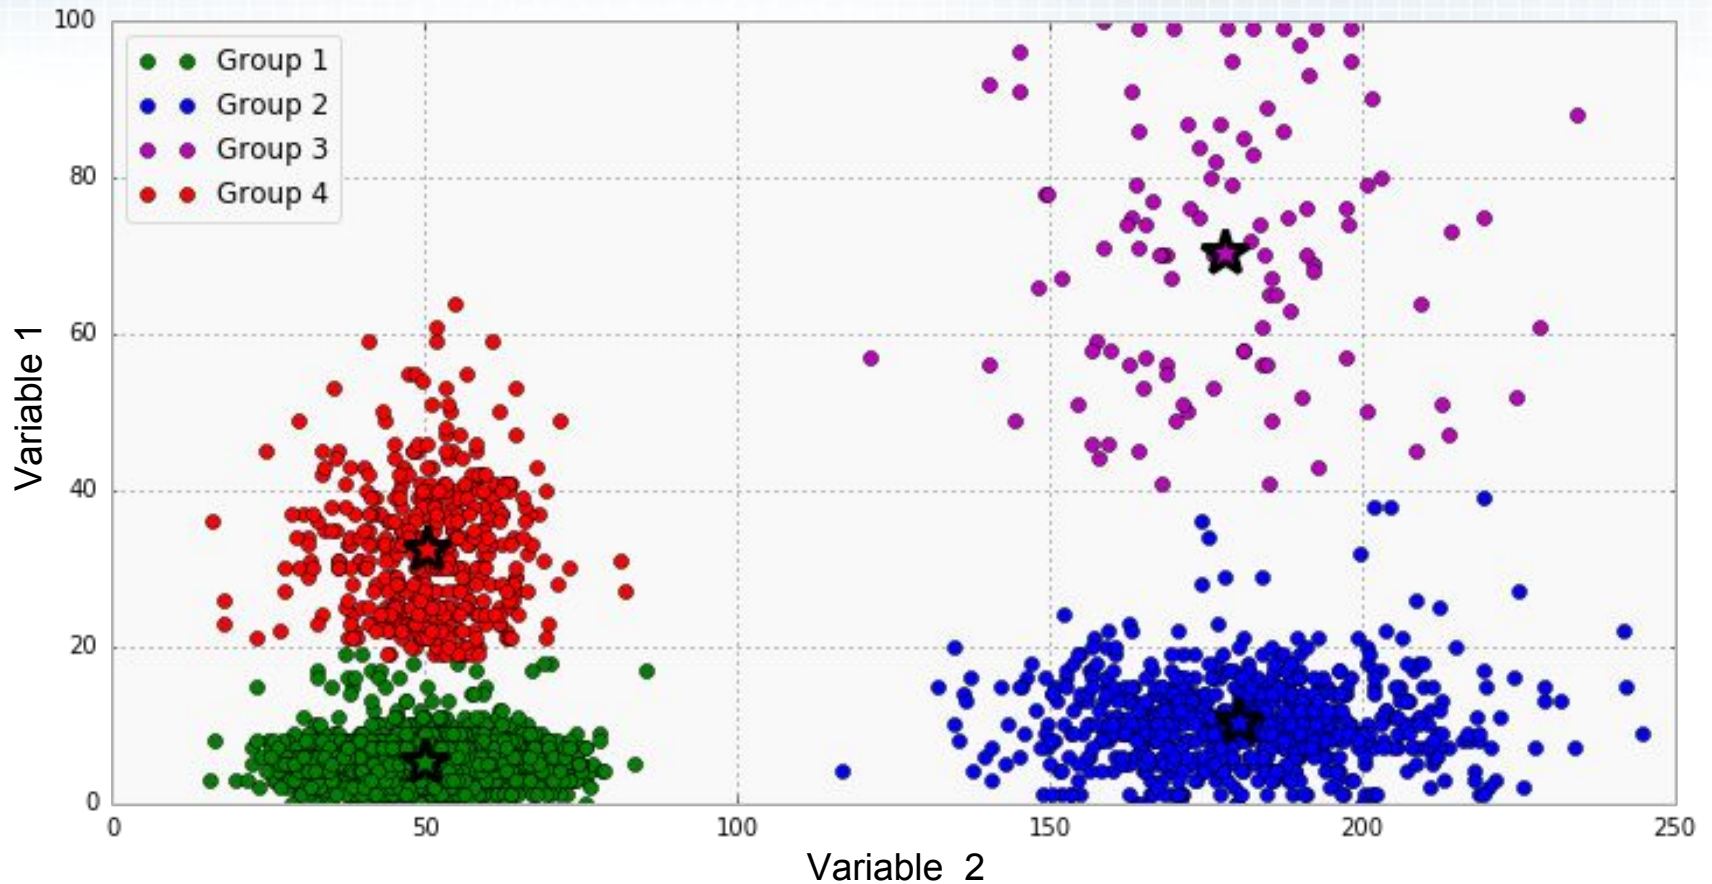
# Course Outline

- **Clusters**
- **Outliers**
- **Regression**
- **Text Analytics**
- **Q & A**

# Clustering

Grouping of similar objects in multivariate data set

Example: grouping by distance from a center



# Clustering with k-means()

- **kmeans()** clustering is the most commonly used algorithm for partitioning a given data set into a set of k groups/clusters
  - k represents the number of groups pre-specified by the analyst. It classifies objects in multiple groups (i.e., clusters), such that objects within the same cluster are as similar as possible
  - In k-means clustering, each cluster is represented by its center (i.e, centroid) which corresponds to the mean of points assigned to the cluster

```
kmeans(x, centers, iter.max = 10, nstart = 1,  
       algorithm = c("Hartigan-Wong", "Lloyd", "Forgy",  
                     "MacQueen"), trace=FALSE)
```

- **nstart()** is optional parameter in kmeans() setting initial number of configurations, e.g. nstart=25 will generate 25 initial random centroids and choose the best one for the algorithm
- **set.seed(seed)** is random number generator, which is useful for creating simulations or random objects that can be reproduced, required by nstart()




# R Script for k-means

```
library(ggplot2)
ggplot(iris, aes(Petal.Length, Petal.Width, color = Species)) + geom_point()
set.seed(20)
irisCluster <- kmeans(iris[, 3:4], 3, nstart = 20)
irisCluster
table(irisCluster$cluster, iris$Species)
irisCluster$cluster <- as.factor(irisCluster$cluster)
ggplot(iris, aes(Petal.Length, Petal.Width, color = irisCluster$cluster)) + geom_point()
```

Or:


```
ggplot(iris, aes(y = Petal.Length, x = seq(1, length(iris$Sepal.Length)), color =
irisCluster$cluster)) + geom_point()
irisCluster1 = kmeans(iris[,3], 3, nstart = 20)
irisCluster1
table(irisCluster1$cluster, iris$Species)
```

# Outliers with Boxplots




outliers

suspected outliers



# Outliers in Iris Boxplot

```
boxplot(Sepal.Length ~ Species, data=iris, col= c("yellow", "blue", "red"),  
ylab="Sepal Length")
```



Why this sample is an outlier?


Lets explore boxplot defaults:

```
boxplot.default(x, ..., range = 1.5, width = NULL, varwidth = FALSE, notch =  
FALSE, names, data = sys.frame(sys.parent()), plot = TRUE, border =  
par("fg"), col = NULL, log = "", pars = NULL)
```


# Outliers in Boxplot

## by Changing Range

```
boxplot(Sepal.Length ~ Species, data=iris, range = 0,  
col= c("yellow", "blue", "red"), ylab="Sepal Length")
```




```
boxplot(Sepal.Length ~ Species, data=iris, range = 1,  
col= c("yellow", "blue", "red"), ylab="Sepal.Length")
```






# Regression with Scatter Plot

```
scatter.smooth(x = cars$speed, y = cars$dist, lpars = list(col = "red", lwd = 3, lty = 3))
```



OR with linear model:

```
linearMod = lm(dist ~ speed, data=cars)  
print(linearMod)
```



Coefficients:

(Intercept)	speed
-17.579	3.932

dist = Intercept + ( $\beta$  \* speed)  
=> dist = -17.579 + 3.932\*speed

# Text Mining

- Natural languages are different from programming languages
  - The semantic or the meaning of a statement depends on the context, tone and other factors
  - Unlike programming languages, natural languages are ambiguous
- 
- Text mining deals with helping computers understand the “meaning” of the text.
  - Some of the common text mining applications include sentiment analysis e.g if a tweet about a movie says something positive or not, text classification e.g classifying the mails you get as spam or ham etc.

# Text Preprocessing

- Text data contains white spaces, punctuations, stopwords etc.
- These characters do not convey much information and are hard to process.
- For example, English stopwords like “the”, “is” etc. do not tell you much information about the sentiment of the text, entities mentioned in the text, or relationships between those entities.
- Depending upon the task at hand, we deal with such characters differently.
- Preprocessing includes
  - Convert the text to lower case, so that words like “write” and “Write” are considered the same word for analysis
  - Remove numbers
  - Remove English stopwords e.g “the”, “is”, “of”, etc
  - Remove punctuation e.g “,”, “?”, etc
  - Eliminate extra white spaces
  - Stemming our text

# Text Processing Vocabulary

- NLP – natural language processing
- Corpus (pl. corpora) – large structured set of texts
- Stopwords – words (such as “the”, “is”, “etc”) that are ignored by search engines
- Stemming - reduce “inflected”/derived word to their stem, e.g. “cars” to “car”
- Lemmatization – distinguish “saw” – hand tool – from “saw” past form of “see”, remove “inflectional” endings to return the base of a word known as lemma
- N-grams – continuous sequence of n items from text, e.g.: phonemes, syllables, letters, words, base pairs
- Tokenization – breaking of texts (strings) into pieces (words, sentences, phases, symbols - tokens) for further analysis.

# Document Term Matrix

- DTM is a matrix that lists all occurrences of words in the corpus, by document
- The documents are represented by rows and the terms (or words) by columns
- If a word occurs in a particular document, then the matrix entry for corresponding to that row and column is 1 (2 if twice, and so on), else it is 0
- Example: Assume we have a simple corpus consisting of two documents, Doc1 and Doc2, with the following content:
  - *Doc1*: bananas are yellow
  - *Doc2*: bananas are good

The DTM for this corpus would look like:

	<i>bananas</i>	<i>are</i>	<i>yellow</i>	<i>good</i>
<i>Doc1</i>	1	1	1	0
<i>Doc2</i>	1	1	0	1

- If we transpose rows and columns we will get TDM



# Text Analytics Steps

- **Create Corpus**
- **Preprocess**
- **Optional: Lemmatization**
- **Create TDM (or DTM) data frame**
  - Rows - document number
  - Columns – words
  - Cells – frequency of a word in a doc
- **Generate 1-gram, bigram and trigram matrices**
- **Analyze**
  - Create summaries
  - Compute Statistics
  - Visualize
    - Histograms
    - Word clouds

# Sentiment Analysis Steps

- **Create Corpus**
- **Preprocess**
- **Stem text**
- **Optional: Lemmatization**
- **Create a subset (training data)**
  - **Manually assign sentiment (e.g. positive Vs. negative (or use sentiment vocabulary)**
  - **Create sentiment score table:**
    - **Rows - document number**
    - **Columns – words**
    - **Cells – score(s) for each word: how often it shows in negative/positive documents**
  - **Test the scores on the test data to obtain confidence level**
- **Apply sentiment scores to all documents for overall results**

# Text Analytics: Wordcloud

```
# i have a dream wordcloud

# Load
library("tm")
library("SnowballC")
library("wordcloud")
library("RColorBrewer")

# Read the text file from internet
filePath <- "http://www.sthda.com/sthda/RDoc/example-files/martin-luther-king-i-have-a-dream-speech.txt"
text <- readLines(filePath)
# Load the data as a corpus
docs <- Corpus(VectorSource(text))
inspect(docs)

# text transformation replacing "/", "@" and "|" with space:
toSpace <- content_transformer(function(x, pattern) gsub(pattern, " ", x))
docs <- tm_map(docs, toSpace, "/")
docs <- tm_map(docs, toSpace, "@")
docs <- tm_map(docs, toSpace, "\\")

# text cleaning
# Convert the text to lower case
docs <- tm_map(docs, content_transformer(tolower))
# Remove numbers
docs <- tm_map(docs, removeNumbers)
# Remove english common stopwords
docs <- tm_map(docs, removeWords, stopwords("english"))
# Remove your own stop word
# specify your stopwords as a character vector
docs <- tm_map(docs, removeWords, c("blabla1", "blabla2"))
# Remove punctuations
docs <- tm_map(docs, removePunctuation)
# Eliminate extra white spaces
docs <- tm_map(docs, stripWhitespace)
# Text stemming
docs <- tm_map(docs, stemDocument)

# build dtm
dtm <- TermDocumentMatrix(docs)
m <- as.matrix(dtm)
v <- sort(rowSums(m), decreasing=TRUE)
d <- data.frame(word = names(v), freq=v)
head(d, 10)

# generate wordcloud
set.seed(1234)
wordcloud(words = d$word, freq = d$freq, min.freq = 1,
  max.words=200, random.order=FALSE, rot.per=0.35,
  colors=brewer.pal(8, "Dark2"))
```

# R References

- <https://www.tidytextmining.com/tidytext.html>
- <http://garrettgman.github.io/tidying/>
- <https://www.r-bloggers.com/intro-to-text-analysis-with-r/>
- <https://eight2late.wordpress.com/2015/05/27/a-gentle-introduction-to-text-mining-using-r/>
  - Required packages: tm, SnowballC, ggplot2, wordcloud
- Tutorial: <https://www.springboard.com/blog/text-mining-in-r/> w/ Hillary's emails
  - Packages for tutorial
    - RSQLite, 'SQLite' Interface for R
    - tm, framework for text mining applications
    - SnowballC, text stemming library
    - Wordcloud, for making wordcloud visualizations
    - Syuzhet, text sentiment analysis
    - ggplot2, one of the best data visualization libraries
    - quanteda, N-grams
  - Datasource: <https://www.kaggle.com/kaggle/hillary-clinton-emails>
- Quanteda
  - <https://cran.r-project.org/web/packages/quanteda/vignettes/quickstart.html>

# Q & A